



# HUST

**ĐẠI HỌC BÁCH KHOA HÀ NỘI**  
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

ONE LOVE. ONE FUTURE.



# GRADUATION THESIS



ĐẠI HỌC  
BÁCH KHOA HÀ NỘI  
HANOI UNIVERSITY  
OF SCIENCE AND TECHNOLOGY

# PDF-TO-DOCX CONVERSION AND PDF EDITING WEB-BASED APPLICATION

- MOTIVATION
- PROBLEM
- PROPOSED SOLUTION
- CONVERSION PIPELINE
- IMPLEMENTATION
- EXPERIMENTAL RESULTS
- CONTRIBUTIONS
- CONCLUSION

ONE LOVE. ONE FUTURE.


- PDF is widely used
- PDF is not editable
- Existing tools have limitations
- Accurate conversion remains challenging



## Why PDF-to-DOCX is difficult?

- No paragraph information
- No reading order
- No semantic labels
- Absolute coordinates
- Complex structured elements


PDF Page



PDF Internal Structure


Raw text "HELLO WORLD"  
Bbox = (121, 223, 497, 355)  
Font = Times New Roman  
Size = 12  
Color = #000000

# PROPOSED SOLUTION




Core idea: Hybrid ML + Rule-based Heuristics

# SYSTEM ARCHITECTURE



- Frontend, backend: Rest APIs
- Independent modules: PDF-DOCX Conversion, PDF Editing
- Conversion engine: ML + Rule-based heuristics
- Backend manages processing jobs
- Local storage

## FastAPI Backend



### Backend responsibilities:

- REST API requests
- Users, documents, processing jobs
- Execute document processing asynchronously
- Generated files for download




## Main features:

- User authentication
- Document management
- PDF editing
- PDF-to-DOCX conversion
- Job monitoring
- Result preview & Download

## Technologies:

- React
- TypeScript


# PDF-TO-DOCX CONVERSION WORKFLOW



## Main Processing Stages:

- PDF content extraction
- Layout analysis
- Coordinate mapping
- Semantic matching
- Content reconstruction
- DOCX generaion

# PDF CONTENT EXTRACTION



## Why PyMuPDF?

- Direct extraction
- Preserve formatting
- Faster than OCR
- High accuracy for digital PDFs



**Output: Structured PDF objects**




Figure 5: Success rates of each model under different feedback conditions. The “Original” point represents the initial success rate for each prompting strategy (Zero-Shot, CoT, and AoT). The “Repeat,” “Specific,” and “Suggestive” points show the final success rates after models were given up to three additional attempts on previously failed puzzles with the corresponding type of feedback. Each line tracks the performance of a prompting strategy across feedback conditions.

100% of trials, with Gemini-2.5-Pro and Llama 3.1 8B-Instruct looping 92% and 86% of the time, respectively. These results suggest that even when presented only with valid moves and the previous move to help track progress, these models struggle to formulate a goal-oriented path.

GPT-5-mini showed a different kind of limitation: it did not solve any puzzles, but its primary outcome was early termination (68%), meaning it consistently made valid moves until it reached the 50-move limit. Therefore, GPT-5-mini exhibits a different kind of planning deficit: it can make valid moves and track the state space—contrasting with its earlier single-prompt results in Fig. 4—but lacks a coherent long-term strategy, effectively wandering the state space without direction.

## 5 Discussion

### 5.1 What did we learn?

Our experiments reveal that a diverse range of LLMs, differing in producer and size, continue to struggle with simple planning and search-based tasks such as the 8-puzzle, even SOTA models like GPT-5-Thinking, which has been described as possessing PhD-level intelligence [33]. Models consistently fail due to two primary deficits: brittle internal state representations that lead to invalid moves and weak heuristic planning that results in loops. These deficits indicate that models often fail to maintain accurate board state representations or to devise effective long-term strategies. This finding underscores a key limitation: while LLMs increasingly excel on academic benchmarks, they remain challenged by sequential tasks that demand sustained planning and state awareness.

The prevalence of loop failures is especially revealing because the system prompt explicitly instructs models to avoid repeating moves. Despite this explicit constraint, all models frequently violated the instruction, suggesting a fundamental difficulty in adhering to constraints. This deficit is most apparent in the external move validator experiment: even when given a list of valid moves, an explicit instruction not to repeat moves, and their previous move, the models continued to loop. One might argue that the model’s internal heuristic changes between calls, but even then, it could simply select a different valid move when its preferred choice violates the constraints. Instead, it repeats moves, indicating that the planning process relies not only on weak heuristics but also fails to integrate the prompt’s explicit rules.

### 5.2 Effects of Prompting

Our experiments suggest that the utility of advanced prompting strategies is model-dependent, as illustrated in Fig. 4. While Algorithm-of-Thought (AoT) and Chain-of-Thought (CoT) improved performance for GPT-5-Thinking, they degraded performance for GPT-5-mini and Gemini-2.5-Pro. These findings demonstrate that, for this task, the efficacy of a prompting technique is contingent on the model, with no single approach proving universally superior.

8

the gains needed to make these models reliable. Furthermore, even when offloading some of the burden by providing the model with valid moves and its last move, no model solves any puzzles. The tendency for models to be unaware of their mistakes and to proceed with high confidence after undetected errors poses a substantial risk in real-world settings. These poor success rates underscore a fundamental gap between current LLM capabilities and the reliable, stateful reasoning required for genuine autonomy.

## 7 Limitations

Our work has several limitations. First, we examine only a single domain. While the 8-puzzle is informative for tracking a model’s progress and may be representative of more challenging sequential planning tasks, it remains a single task, which limits our ability to draw broad conclusions about model performance in other settings. Second, we evaluate the models on only 50 puzzles, which is a relatively small dataset. Finally, although we generated puzzles at random to reduce the risk of training-data contamination, we cannot entirely rule it out.

## References

- [1] OpenAI. GPT-5 system card. <https://openai.com/index/gpt-5-system-card/>. 2025. Web article, accessed 2025-11-10.
- [2] Gemini Team. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv*, 2025. URL <https://arxiv.org/abs/2507.06261>. arXiv:2507.06261.
- [3] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv*, 2020. URL <https://arxiv.org/abs/2009.03300>. arXiv:2009.03300.
- [4] Mislav Balunović, Jasper Dekoninck, Ivo Petrov, Nikola Jovanović, and Martin Vechev. Matharena: Evaluating LLMs on uncontaminated math competitions. *arXiv*, 2025. URL <https://arxiv.org/abs/2505.23281>. arXiv:2505.23281.
- [5] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv*, 2023. URL <https://arxiv.org/abs/2310.06770>. arXiv:2310.06770.
- [6] Karthik Valmeekam, Kaya Stechly, Atharva Gundawar, and Subbarao Kambhampati. A systematic evaluation of the planning and scheduling abilities of the reasoning model o1. *Transactions on Machine Learning Research*, 2025. URL <https://openreview.net/forum?id=PKKzxp0Pnk>. Reviewed on OpenReview.
- [7] François Chollet, Mike Knoop, Gregory Kamradt, and Bryan Landers. ARC-prize-2024: Technical report. *arXiv*, 2024. URL <https://arxiv.org/abs/2412.04604>. arXiv:2412.04604.
- [8] Marianna Nezhurina, Lucia Cipolina-Kun, Mehdi Cherti, and Jenia Jitsev. Alice in wonderland: Simple tasks showing complete reasoning breakdown in state-of-the-art large language models. *arXiv*, 2024. URL <https://arxiv.org/abs/2406.02061>. arXiv:2406.02061.
- [9] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. Agentbench: Evaluating LLMs as agents. *arXiv*, 2023. doi: 10.48550/arXiv.2308.03688. URL <https://arxiv.org/abs/2308.03688>. arXiv:2308.03688.
- [10] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, 2022.
- [11] Zihan Yu, Liang He, Zhen Wu, Xinyu Dai, and Jiajun Chen. Towards better chain-of-thought prompting strategies: A survey. *arXiv*, 2023. URL <https://arxiv.org/abs/2310.04959>. arXiv:2310.04959.

11

## DocLayout-YOLO

- Detect semantic regions
- Classify regions into categories
- Bouding boxes of detected regions

## 11 classes:

- Title
- Text
- Table
- Figure caption
- Figure
- Table caption
- Others

# COORDINATE MAPPING

## Problem:

Different coordinate systems:

- PyMuPDF: points
- DocLayout-YOLO: pixels



## Solution:

- Convert PDF bounding boxes into image coordinates
- All object in a unified coordinate space

PyMuPDF Coordinate Space  
(Unit: Points)

Hello World

(41, 342, 195, 375)

DocLayout-YOLO Coordinate Space  
(Unit: Pixels)

Hello World

(55, 455, 260, 500)

$$1 \text{ pt} = 1.33\text{px}$$

# SEMANTIC MATCHING

## PyMuPDF



## DocLayout-YOLO



→ Title

## Matching strategy:

- Compare mapped elements with detected layout regions
- Measure overlap = IoU
- Assign semantic label to best matching

$$IoU = \frac{\text{Intersection Area}}{\text{Union Area}}$$

**Semantic Blocks**



**Reading Order**



**Paragraphs**



**Formatting**



**DOCX**

## Reconstruction Process:

- Recover logical reading order
- Merge related text spans into paragraphs
- Preserve formatting
- Convert document elements into DOCX objects

## Structural Refinement:


- Tight bounding box refinement
- Paragraph grouping
- Cross-page section continuity
- IoU & Containment ratio elements tracking

## Content Refinement:

- External hyperlinks preservation
- Text properties preservation



## Paragraph grouping



## Tight bounding box



## Cross-page section continuity

### 6 Conclusion

In this work, we examined a range of LLMs using the 8-puzzle to probe their planning and stateful reasoning capabilities. Our findings show that both SOTA and small models still struggle significantly with this seemingly simple task, with failures indicating critical deficits in state tracking and heuristic planning. While feedback and prompt engineering offer some help, they are costly and do not provide

10

Adobe Acrobat

AutoDoc

the gains needed to make these models reliable. Furthermore, even when offloading some of the burden by providing the model with valid moves and its last move, no model solves any puzzles. The tendency for models to be unaware of their mistakes and to proceed with high confidence after undetected errors poses a substantial risk in real-world settings. These poor success rates underscore a fundamental gap between current LLM capabilities and the reliable, stateful reasoning required for genuine autonomy.

### 7 Limitations

Our work has several limitations. First, we examine only a single domain. While the 8-puzzle is informative for tracking a model's progress and may be representative of more challenging sequential planning tasks, it remains a single task, which limits our ability to draw broad conclusions about


## IoU & Containment ratio tracking

| Model                 | Approach | Feedback   | Solved N | Time (min) | Tokens | Attempts | Moves | Optimal |
|-----------------------|----------|------------|----------|------------|--------|----------|-------|---------|
| GPT-5-Thinking        | AoT      | Suggestive | 34       | 23.92      | 75284  | 1.97     | 48.91 | 21.03   |
| Gemini-2.5-Pro        | Base     | Repeat     | 9        | 6.74       | 58651  | 2.67     | 30.56 | 19.67   |
| Llama 3.1 8B-Instruct | Base     | Repeat     | 0        | 0          | 0      | 0        | 0     | 0       |
| GPT-5-mini            | CoT      | Repeat     | 9        | 9.91       | 44272  | 2.67     | 28.44 | 19.11   |


| Model                 | Approach | Feedback   | Solved N | Time (min) | Tokens | Attempts | Moves | Optimal |
|-----------------------|----------|------------|----------|------------|--------|----------|-------|---------|
| GPT-5-Thinking        | AoT      | Suggestive | 34       | 23.92      | 75284  | 1.97     | 48.91 | 21.03   |
| Gemini-2.5-Pro        | Base     | Repeat     | 9        | 6.74       | 58651  | 2.67     | 30.56 | 19.67   |
| Llama 3.1 8B-Instruct | Base     | Repeat     | 0        | 0          | 0      | 0        | 0     | 0       |
| GPT-5-mini            | CoT      | Repeat     | 9        | 9.91       | 44272  | 2.67     | 28.44 | 19.11   |

Table 1: Solved count and average time (minutes), tokens used, attempts, moves, and optimal moves for each model under its best performing condition. We broke ties for GPT-5-mini and Gemini-2.5-Pro by choosing the approach that required the fewest tokens.

## External hyperlinks preservation



## Text properties preservation



Structured Elements



Paragraph




Run



Image



DOCX File



**Editable Microsoft Word Document**

# PDF EDITING FEATURES


Merge PDF



Split PDF




Organize PDF



Insert pages



Number pages



Reorder pages



Crop pages





DEMO

# CONVERSION RESULT COMPARISON

| Feature             | AutoDoc | Adobe Acrobat | iLovePDF |
|---------------------|---------|---------------|----------|
| Layout              | ✓       | ✓             | ✓        |
| Font properties     | ✓       | ✓             | ✓        |
| Images              | ✓       | ✓             | ✓        |
| Editable tables     | ✗       | ✓             | ✓        |
| Formulas            | ✗       | ✗             | ✗        |
| Section continuity  | ✓       | ✗             | ✗        |
| Internal hyperlinks | ✗       | ✓             | ✓        |
| External hyperlinks | ✓       | ✓             | ✓        |
| Multi-column        | ✓       | ✓             | ✓        |
| Speed               | ✗       | ✓             | ✓        |

## Hybrid PDF-to-DOCX Pipeline

- DocLayout-YOLO + Rule-based reconstruction

## Semantic Document Reconstruction

- Logical document structure
- Section continuity across page boundaries

## Reconstruction Refinement

- Improve layout quality
- Preserve hyperlinks, text format

## Practical Conversion Quality

- Editable DOCX with high layout fidelity
- Comparable results with commercial tools



## Limitations

- “Text swallowing” due to AI imprecision
- Pagination discrepancy
- Unprocessed internal hyperlinks
- Section header without numeric prefixes
- Non-editable complex elements

## Future work

- Support internal hyperlinks
- Optimize processing performance
- Reconstruct editable tables, formulas

## Project summary

- AutoDoc, a web-based PDF processing system
- Hybrid PDF-to-DOCX conversion pipeline, ML + rule-based heuristics
- PDF editing and document conversion
- Preserve document layout, formatting, hyperllinks
- Comaprable conversion quality with commercial tools



A large graphic on the left side of the slide. It features a dark blue background with a circular pattern of red dots of varying sizes, creating a sense of depth and movement. The word "HUST" is centered within this pattern in a white, bold, sans-serif font.

# HUST

**Thank you  
for your attention!**